

Machine Learning Adversarial Attacks

Internship Evaluation

Understanding Security Vulnerabilities in AI Systems

By Parul Gupta
12901172023
CSE AI 2

Contents

01. Introduction

Overview of adversarial attacks and their impact on ML systems

02. Attack Types

Classification of adversarial attack methods and techniques

03. Attack Mechanisms

How adversarial examples are generated and optimized

04. Defense Strategies

Current approaches to mitigate adversarial vulnerabilities

What Are Adversarial Attacks?

Core Concept

Adversarial attacks are deliberate perturbations added to input data that cause machine learning models to make incorrect predictions while appearing unchanged to humans.



Stealth Nature

Imperceptible changes that fool AI systems



Security Risk

Critical vulnerabilities in deployed models



Attack Classification Matrix

White-box Attacks

Full model knowledge: gradients, architecture, parameters

S

W

Black-box Attacks

Limited knowledge: query-based, transfer attacks

Non-targeted Attacks

Any incorrect prediction suffices


T

O

Targeted Attacks

Specific misclassification goals

Fast Gradient Sign Method (FGSM)



Gradient Computation

Calculate loss gradient with respect to input data



Perturbation Generation

Apply sign of gradient with small epsilon



Adversarial Creation

Add perturbation to original input


Attack Success Rates by Model Type

Vulnerability Analysis

Deep neural networks show varying susceptibility to adversarial attacks across different architectures and training methods.

CNNs: 85-95% attack success rate,
RNNs: 70-80%, Transformers: 60-75%

Attack Success Distribution



Attack Generation Pipeline

Input Selection

Choose clean sample from target dataset



Gradient Analysis

Compute model's sensitivity to input perturbations



Perturbation Crafting

Generate minimal adversarial noise using optimization




Validation Testing

Verify attack effectiveness and imperceptibility

Real-World Attack Scenarios

Attack Impact by Domain



Autonomous Vehicles

Stop sign
misclassification
on causing
safety risks




Medical Imaging

Cancer
diagnosis
manipulation
with subtle
pixel changes

Defense Strategy Framework

Adversarial Training

Incorporate adversarial examples during model training to improve robustness against known attack patterns



Input Transformations

Apply preprocessing techniques like denoising, compression, or randomization to reduce perturbation effects

Certified Defenses

Mathematical guarantees of robustness within specified perturbation bounds using formal verification methods

Defense Effectiveness Metrics

85%

Adversarial Training

72%

Input Denoising

94%

Ensemble Methods

91%

**Certified
Robustness**

Adversarial Training

Most practical defense method

- Retrains on adversarial examples
- Improves robustness significantly
- Computationally expensive

Certified Defenses

Strongest theoretical guarantees

- Provides mathematical guarantees
- Limited to small perturbations
- High computational complexity

Detection Methods

Complementary to robust training

- Identifies adversarial inputs
- Rejects suspicious samples
- May increase false positives

Future Research Directions

Emerging Challenges

Next-generation adversarial attacks and defenses require addressing more sophisticated threat models and real-world deployment constraints



Physical Attacks

Adversarial patches in real-world environments



Universal Perturbations

Single perturbation across multiple inputs



Adaptive Attacks

Attacks that evolve against new defenses



Thank You

Questions and Discussion Welcome